



TAIS NextGen-Tunisian AI Society

SympactAI



Hands-On Reinforcement Learning For Stock Trading Using FinRL

Supervised by

Mr.Mahdi Kallel

Mrs.Imen Ghazouani

Presented by

Moez KESSEMTINI

Mazen TMANI

Saturday, August 16, 2025

Outline

01

Opening Hook

02

**RL &
Stock Trading**

03

**FinRL & Trading
Environment**

04

**RL Workflow in
FinRL**

05

**Results &
Evaluation**

06

Future Directions

01

Opening Hook



Problematic



About 95% of retail traders lose money due to behavioral biases like overconfidence and poor risk management (Medium 2022)

Jane Street (India, 2023–2025): Retail investors lost \$21.7B in derivatives trading (Reuters, 2025)

**Can algorithms and AI solve those
Problems ?**

Motivation



- In 2023, **37%** of overall U.S. equity trading volume was executed through algorithms
- The global reinforcement learning market is expected to grow from \$2.8B in 2022 to **\$88.7B** by 2032

What about implementing RL algorithms in trading ?

02

RL & Stock Trading



Core Concepts of Reinforcement Learning

Agent: the learner or decision-maker

Environment: the system or market the agent interacts with

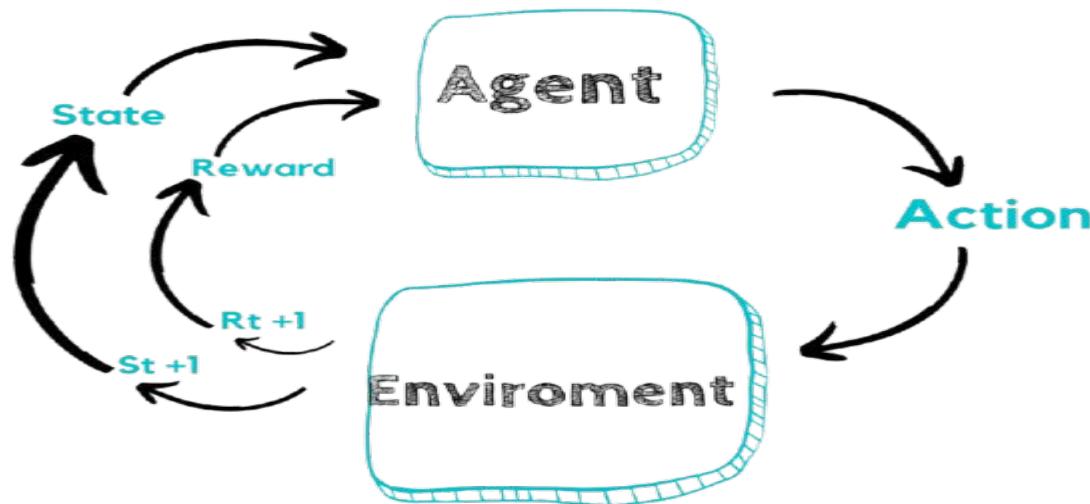
State: current situation or information available to the agent

Action: decisions or moves taken by the agent

Reward: feedback received based on the action's outcome

Policy: strategy that maps states to actions

Value Function: estimates expected future rewards





Why RL is Ideal for Stock Trading



Dynamic & Uncertain Markets

RL agents learn by trial and error, adapting to changing market conditions

Strategies continuously evolve by learning from new market data and adapting to emerging trends

Advantages over Traditional Methods:

Adapts in real-time

Handles complex, nonlinear relationships

Incorporates risk management constraints (drawdown limits, portfolio volatility)

03

FinRL & Trading Environment



What is FinRL?



Open-source RL framework developed by
AI4Finance



Built on: PyTorch, Stable Baselines3, Gym interface



Pre-built trading environments



Integration with financial data sources
(Yahoo Finance, Quandl...)



ideal for experimentation, research,
and practical RL trading projects



FinRL Trading Environment

Components

- **State space:** prices, technical indicators, portfolio info
- **Action space:** buy, sell, hold (and position sizing)
- **Reward function:** portfolio value, Sharpe ratio

Realistic constraints

- Transaction costs
- Risk management parameters

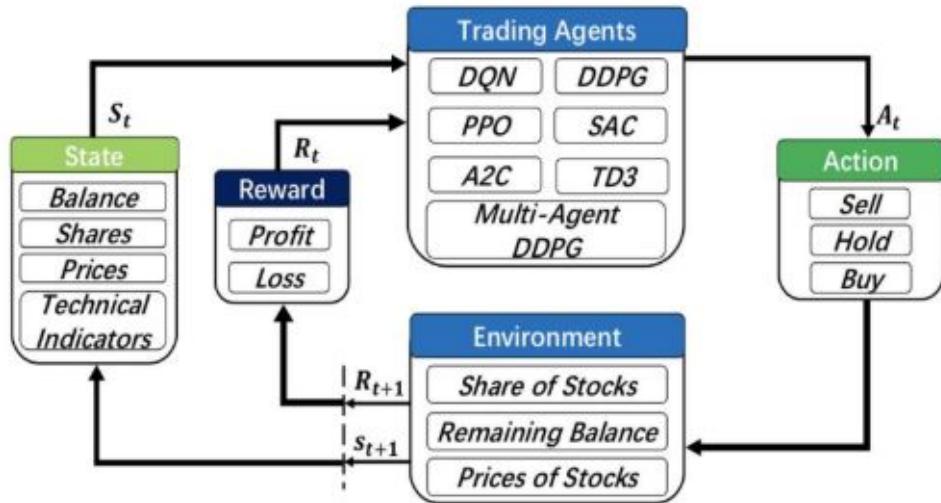
Automated Trading in FinRL

Agent:

- Represent the investor/trader.
- Perceive the current state in the market.
- Take an action for trading.
- Get rewards from the action.
- Learn from past trading experience.

Environment:

- Represent the trading market.
- State is a snapshot of current market conditions.
- Receive the trading action and step into a new state.
- Return the reward for the trading action.



Overview of automated trading in FinRL,
using deep reinforcement learning.

FinRL Layers :

Applications

Stock Trading,
Portfolio Allocation

High-Frequency
Trading

Cryptocurrency
Trading

Market
Regulations

User-defined
Tasks

DRL Agents

DRL Libraries: *ElegantRL, RLlib, Stable Baselines 3*

DQN, D3QN
Double DQN

DDPG

TD3

A2C,
SAC

PPO

MADDPG,
MAPPO

User-designed
DRL Algorithms



Market Environments

Historical Data API:
WRDS
Yahoo! Finance

Live Trading API:
CCXT, Alpaca,
QuantConnect

Market
Simulations

User-imported
Datasets

04

RL Workflow in FinRL



Data Layer

Environment Layer

Agent Layer

Unified Data Processor



State-Action-Reward

Plug-and-Play DRL Agents/Libraries



*Paper trading
or live trading*

Adjust
factors

Adjust
environment settings

Validation

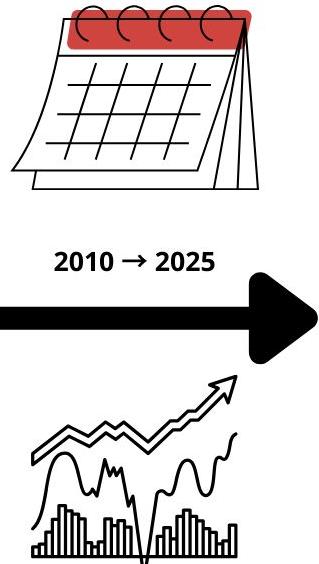
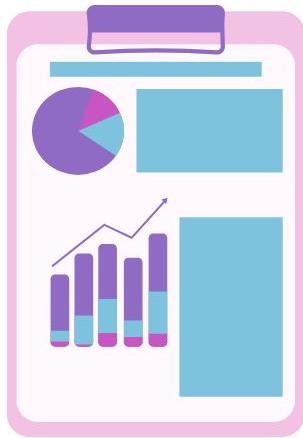
Training

Testing

Trading

Adjust
parameters

Data Collection

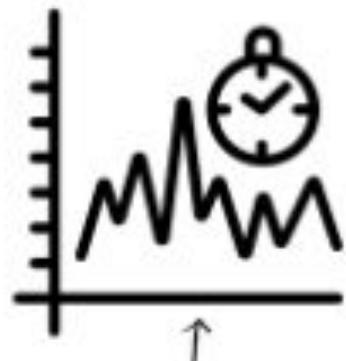
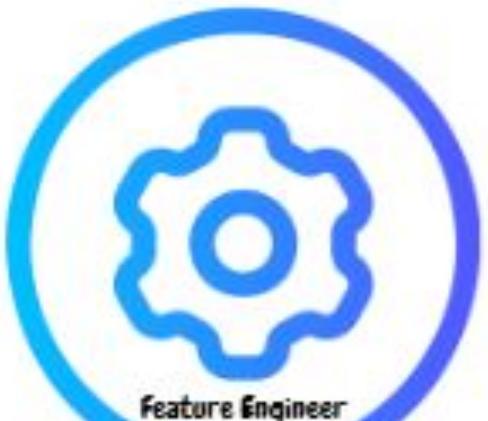


Yahoo Finance
(daily OHLCV data)

30 Dow Jones
Industrial Average stocks

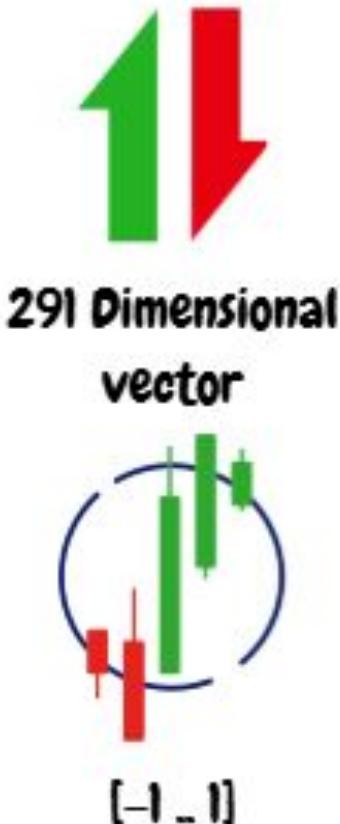
Price	date	close	high	low	open	volume	tic	day
0	2010-01-04	6.424604	6.439314	6.375672	6.407193	493729600	AAPL	0
1	2010-01-04	39.913254	40.016977	39.111116	39.159521	5277400	AMGN	0
2	2010-01-04	32.637966	32.781535	32.215237	32.550232	6894300	AXP	0
3	2010-01-04	43.777550	43.941189	42.702201	43.419101	6186700	BA	0
4	2010-01-04	39.403461	39.834174	38.703553	38.797774	7325600	CAT	0

Data Preprocessing

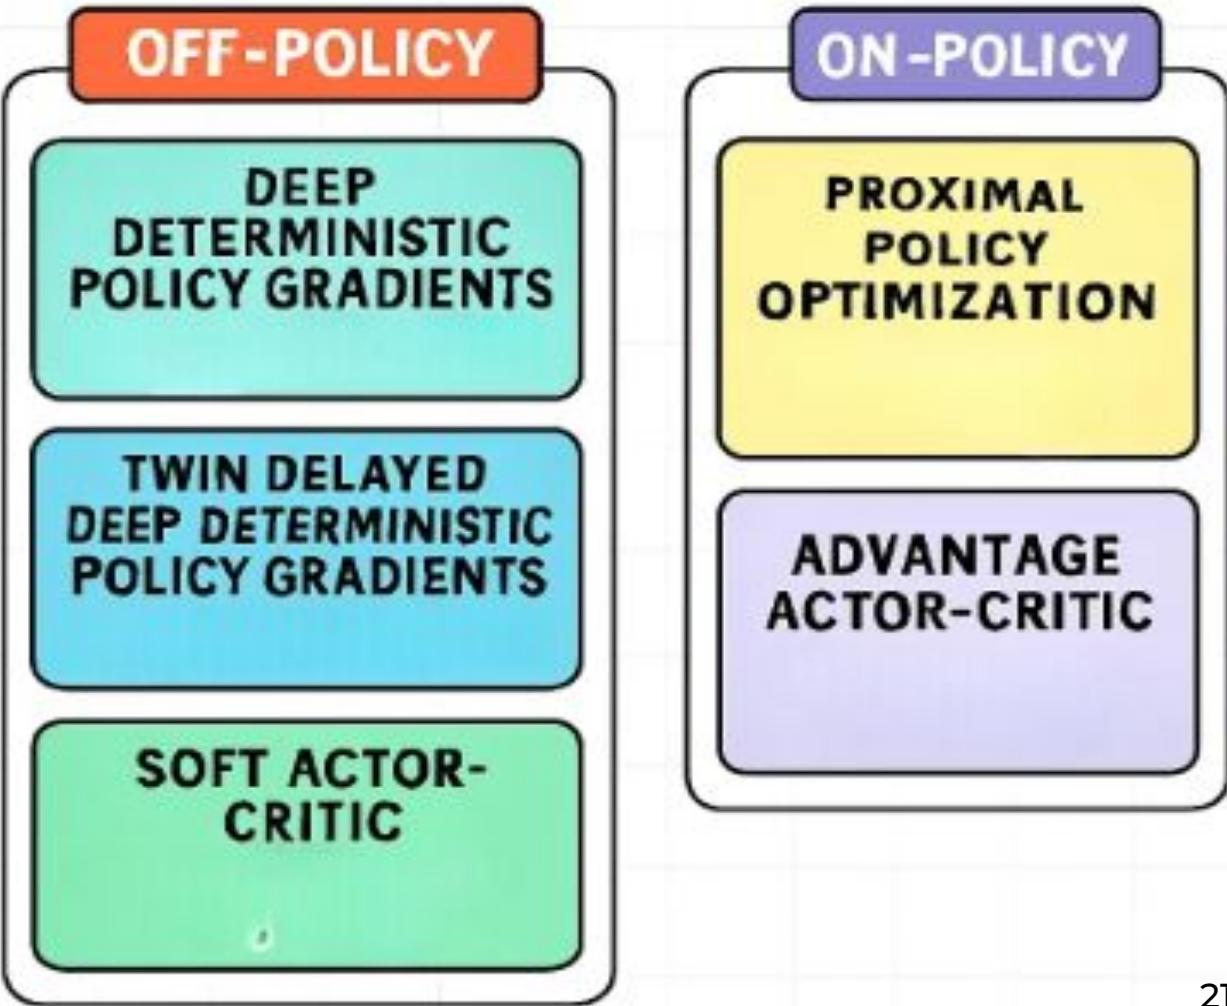


Technical and Market Indicators

Stock Trading Environment



RL Algorithms and Libraries



Library	Pros	Cons
Stable-Baselines3 	Simple to set up, well-documented, large community, stable & reproducible results, widely used in finance	Limited distributed training
ElegantRL	Fast GPU training, flexible for custom algorithms	Smaller community, less documentation
RLlib	Highly scalable, supports massive parallelism	Complex setup, slower for small/ medium tasks

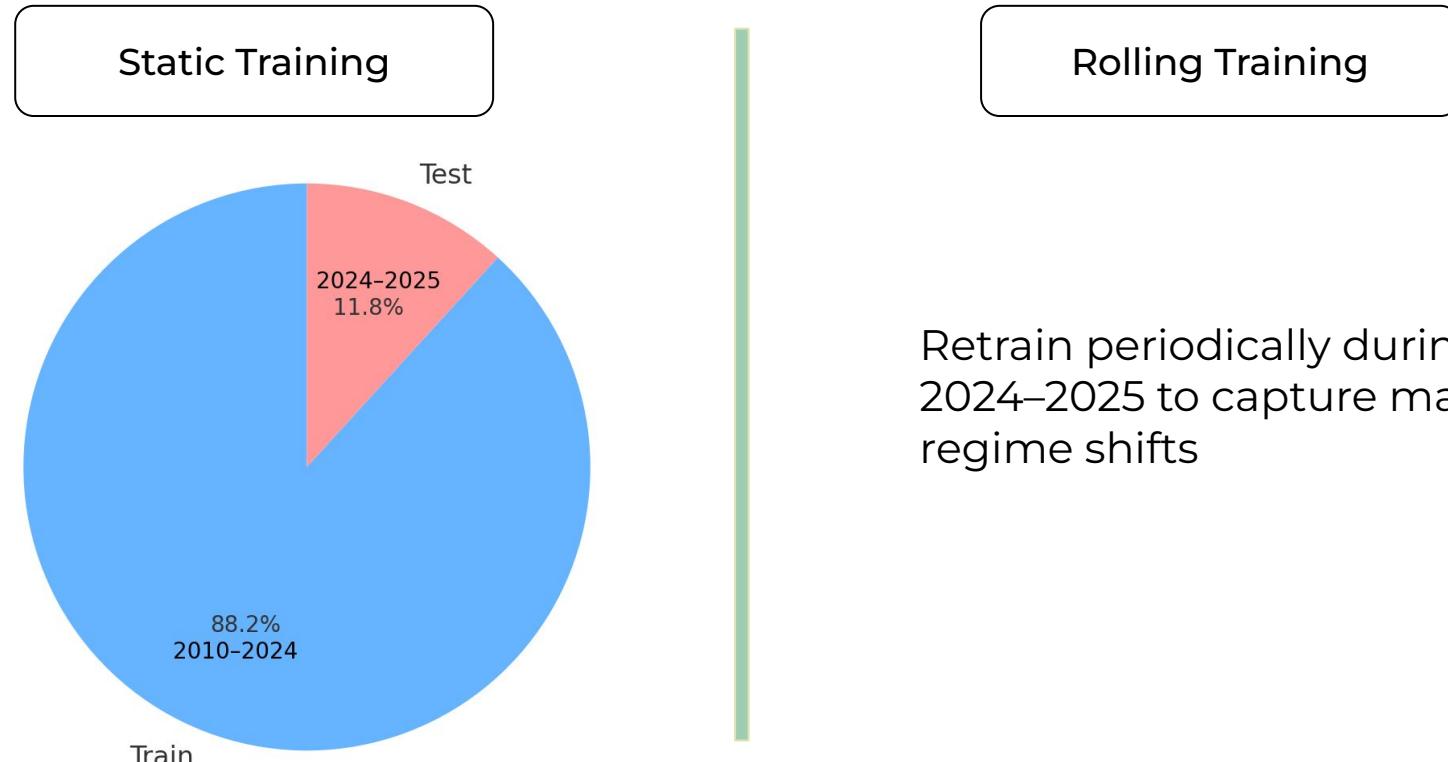
Agents Training

Training Paradigms

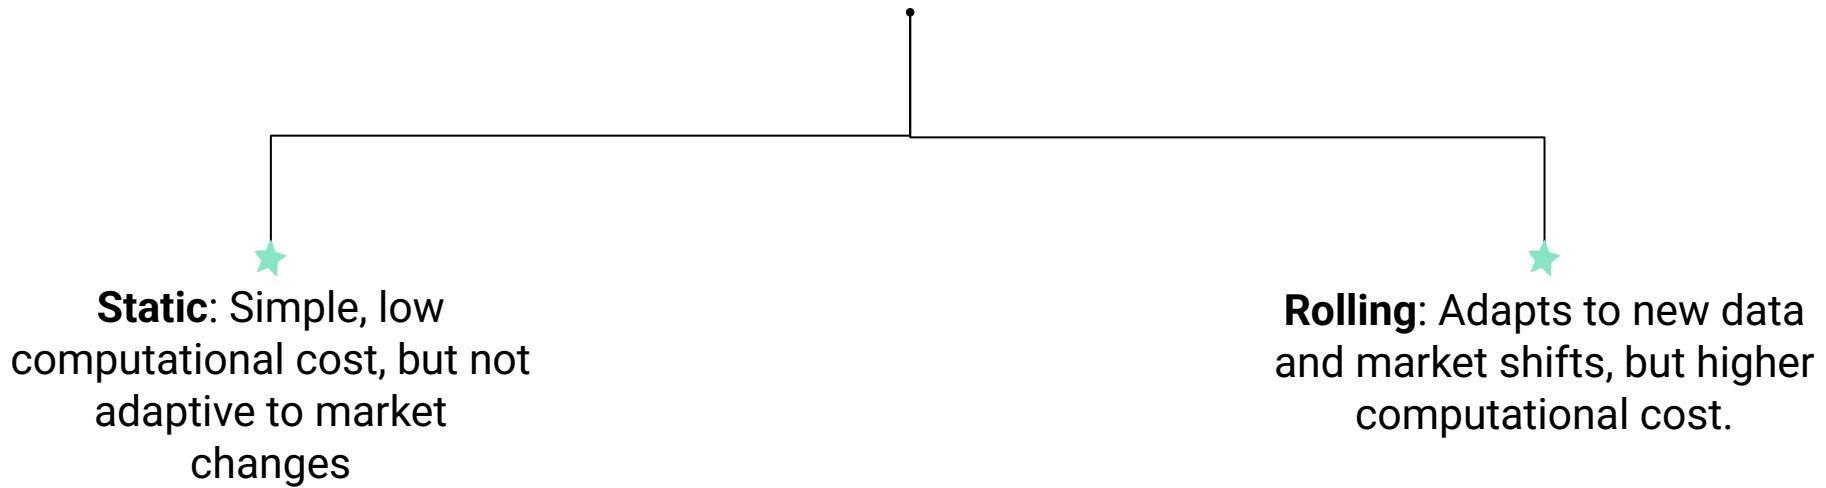
Static Training: Train once on historical data, then test on future unseen data.

Rolling Training: Periodically retrain with new data to adapt to evolving market conditions.

Data Distribution :



Trade-offs

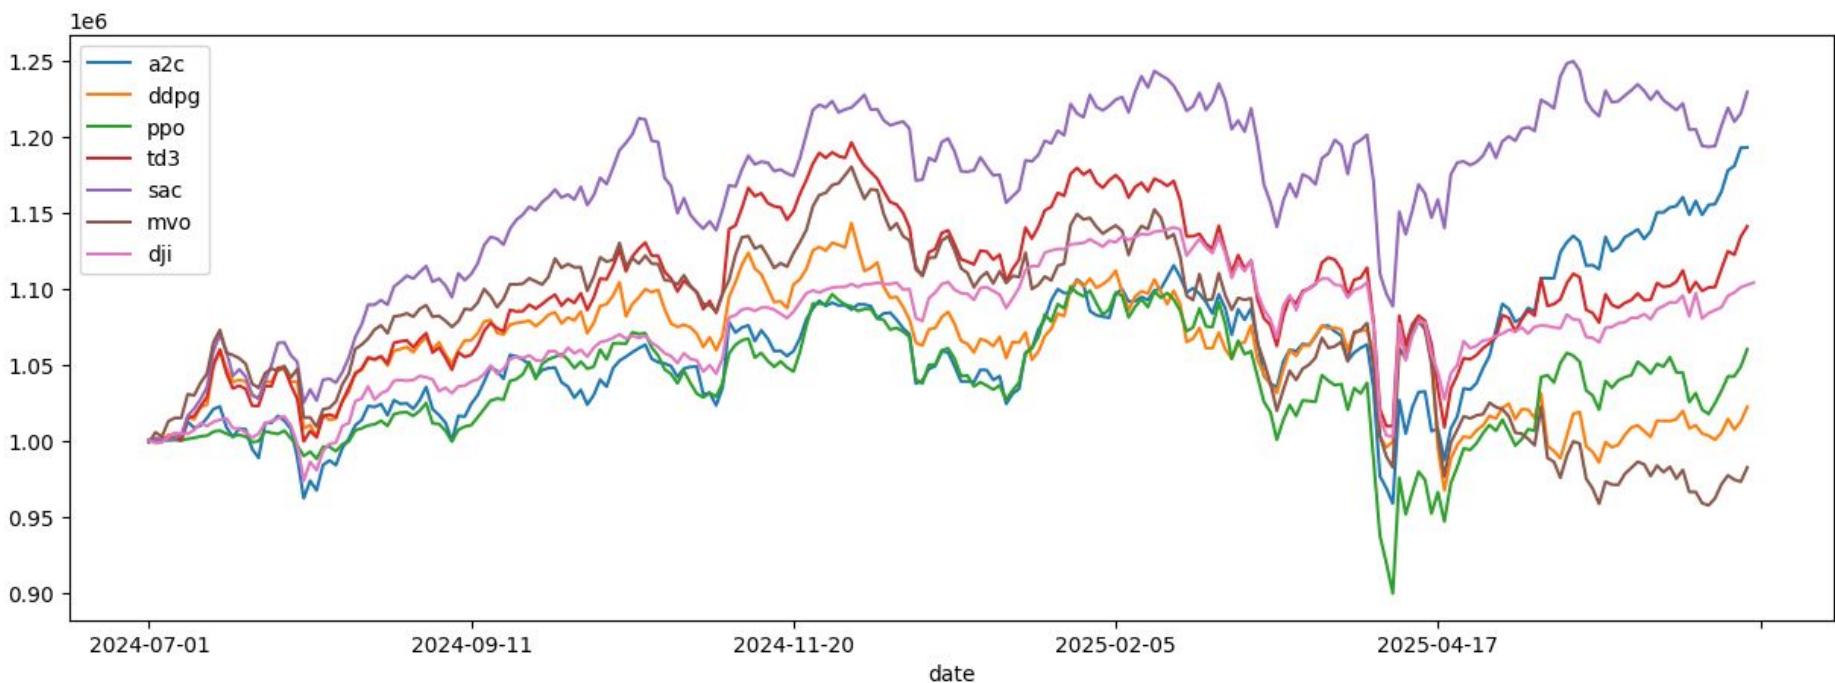


05

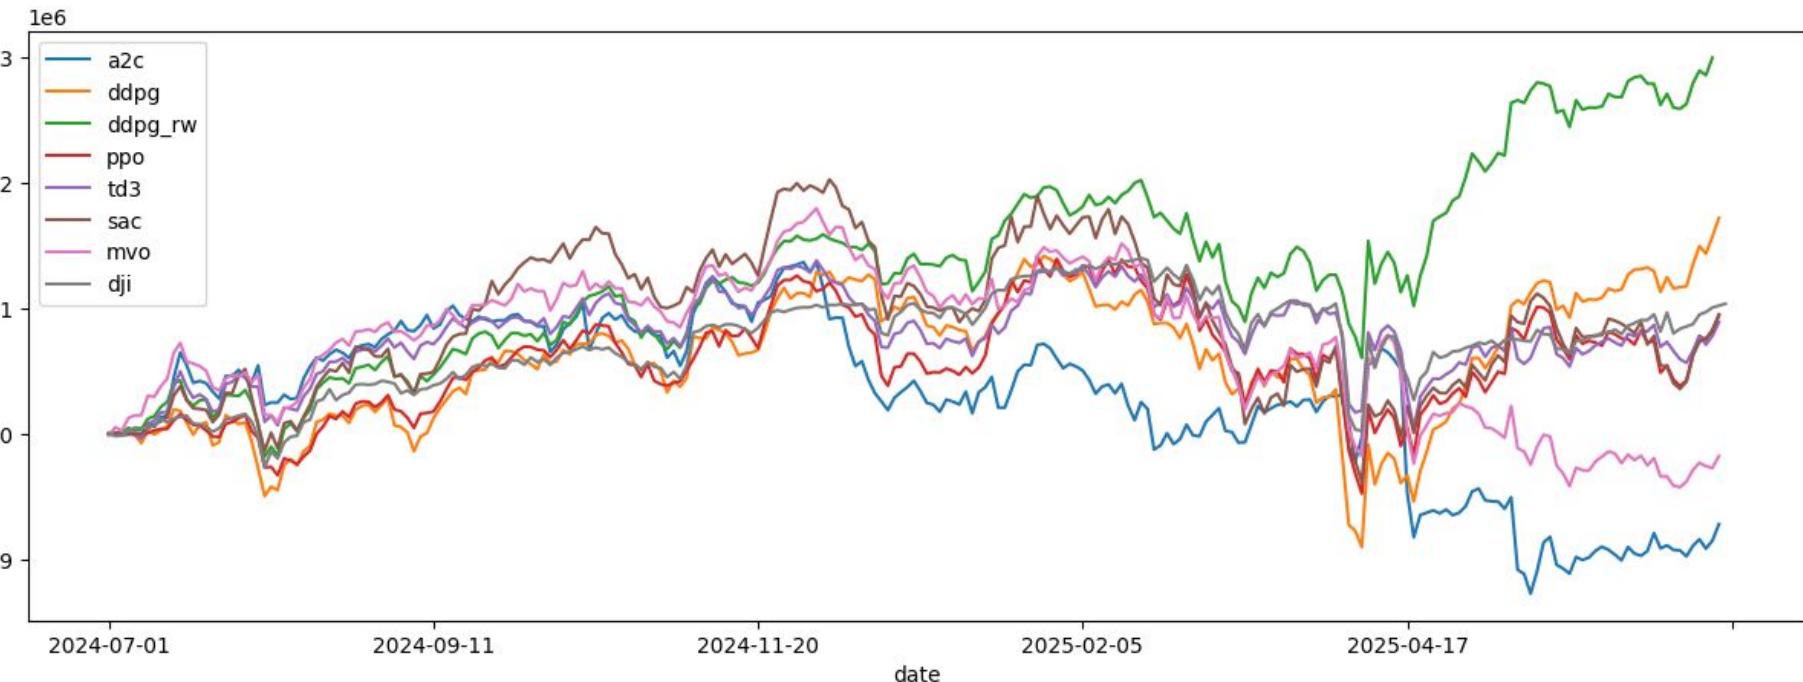
Results & Evaluation



Cumulative returns for DRL agents and benchmarks



Performance with DDPG retrained via expanding rolling window



Selected performance metrics for all Agents

	PPO	DDPG	DDPG	RW	A2C	TD3	SAC
Annual return	0.095	0.174	0.306	-0.073	0.090	0.097	
Sharpe ratio	0.603	0.926	1.649	-0.268	0.640	0.571	
Calmar ratio	0.575	0.858	2.591	-0.313	0.768	0.481	
Max drawdown	-0.166	-0.203	-0.118	-0.233	-0.117	-0.201	
Stability	0.143	0.187	0.793	0.501	0.067	0.003	



Agent performance varies by algorithm and settings



Using rolling window training, DDPG outperforms all other RL algorithms



The rolling window strategy proved effective in boosting returns and reducing risk

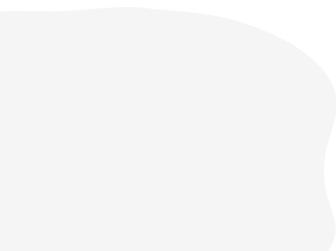
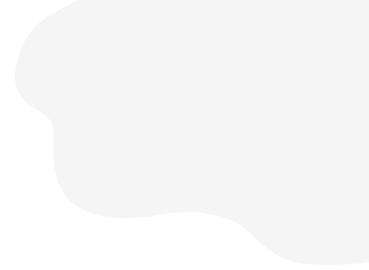
06

Future Directions



- **Integration with News & Sentiment Analysis:** leverage financial news and social media for improved decision-making
- **Real-Time Trading Deployment:** deploy RL agents in live or paper trading environments
- **Ensemble of RL Agents:** combine multiple RL strategies for more robust performance





**Thank you for your
attention**